

## 1. With a neat diagram, explain the concept of the virtual machine.

### Virtual Machine: Concept and Example

#### Definition:

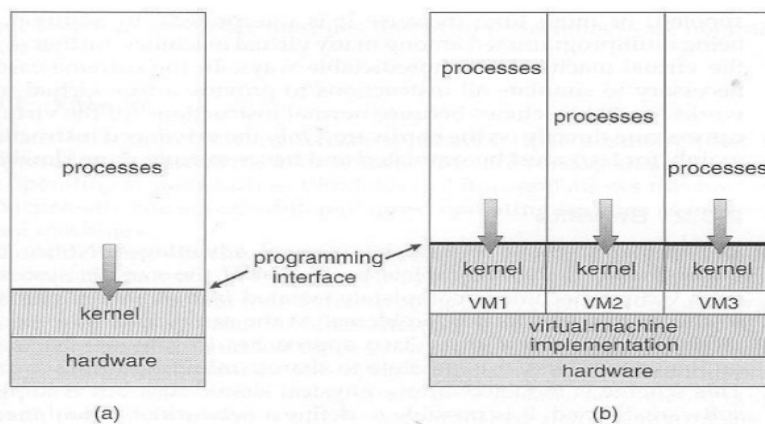
A virtual machine (VM) abstracts the hardware of a computer (CPU, memory, disk, etc.) into multiple independent execution environments, making each environment appear like a private computer to the user.

#### Key Points:

- Each process has the illusion of its own processor and memory.
- The **host OS** is the main operating system installed on the system, while the additional operating systems installed are called **guest OS**.
- VMs were first implemented in the **VM Operating System** for IBM mainframes in 1972.

#### Implementation:

- Virtual machine software runs in **kernel mode** (like an OS), while the virtual machine itself executes in **user mode**.
- Implementation is complex as it requires duplicating the underlying hardware environment.



**Figure: System modes. (A) Non-virtual machine (b) Virtual machine**

#### Benefits of Virtual Machines:

##### 1. Resource Sharing:

- Multiple operating systems can share the same hardware.

##### 2. Isolation:

- Guest OS are isolated from one another and the host OS. For example, a virus in one guest OS cannot affect others or the host system.

##### 3. Software Sharing:

- Resources can be shared via shared file systems or virtual communication networks.

##### 4. Development and Testing:

- Developers can run multiple OS environments simultaneously, allowing rapid code testing and system development.

##### 5. System Consolidation:

- Multiple systems can be combined into one physical system, reducing hardware requirements.

**Example:**

Running **Windows** as the host OS and installing **Ubuntu** and **macOS** as guest OS in a virtualization software like **VMware** or **VirtualBox**. Each OS can run independently, sharing the same physical resources like CPU and memory.

---

## **2. What are system calls? Briefly point out its types with illustrations.**

### **System Calls and Their Types**

**Definition:**

System calls provide an interface between a program and the operating system, enabling the program to request services from the OS. They are often written in C or C++ and sometimes in assembly for performance optimization.

**How System Calls Work:**

- System calls are used in tasks like copying data from one file to another.
  - The program may use system calls to open files, read from the input file, write to the output file, handle errors (e.g., missing files), and terminate the program.
- Most programmers use APIs to access system calls indirectly, ensuring portability across systems.
- System call numbers in a **system call table** determine the appropriate service routine to execute.

**Parameter Passing Methods:**

1. **In registers**
2. **Passing the address of a parameter block** (used in Linux & Solaris)
3. **Using the stack** (parameters pushed by the program and popped by the OS).

**Types of System Calls:**

1. **Process Control**
  - Examples: Create process, terminate process, load/execute program.
2. **File Management**
  - Examples: Open, read, write, close files.
3. **Device Management**
  - Examples: Request/release devices, read/write to devices.
4. **Information Management**
  - Examples: Get/set system data, retrieve file attributes.
5. **Communications**
  - Examples: Send/receive messages, establish/terminate connections.
6. **Protection**
  - Examples: Control access permissions, manage user credentials.

### 3. Explain the states of a process with a transition diagram and process control block.

#### Process States and Process Control Block (PCB)

##### Process States:

A process can exist in the following five states:

1. **New:**
  - The process is being created.
2. **Ready:**
  - The process has all required resources and is waiting to be assigned to the CPU.
3. **Running:**
  - The process is currently executing its instructions.
4. **Waiting:**
  - The process is waiting for an event (e.g., I/O operations, user input) to occur.
5. **Terminated:**
  - The process has completed execution and is removed from the system.




Figure: Diagram of process state

---

#### Process Control Block (PCB):

The **PCB** is a data structure that stores all information related to a specific process. It contains:

1. **Process State:**
  - The current state (new, ready, running, waiting, or terminated).
2. **Program Counter:**
  - Address of the next instruction to be executed by the process.
3. **CPU Registers:**
  - Information like accumulators, stack pointers, and index registers, saved for process resumption.
4. **CPU Scheduling Information:**

- Includes process priority, scheduling parameters, and pointers to scheduling queues.

#### 5. **Memory Management Information:**

- Includes base and limit registers, page tables, or segment tables.

#### 6. **Accounting Information:**

- Tracks CPU usage, time limits, process ID, and other details.

#### 7. **I/O Status Information:**

- List of I/O devices and open files associated with the process.



Figure: Process control block (PCB)

- 
4. Calculate average waiting and turnaround times by drawing the Gantt chart using FCFS and RR( $q=2\text{ms}$ ).

Processes	Arrival Time	Burst Time
P1	0	9
P2	1	4
P3	2	9
P4	3	5

---

### 5. Explain five different scheduling criteria used in the computing scheduling mechanism.

#### Five Scheduling Criteria in CPU Scheduling

##### 1. **CPU Utilization:**

- Measures how effectively the CPU is being used.
- Goal: Keep CPU busy as much as possible (ranges from 40% in light systems to 90% in heavy systems).

##### 2. **Throughput:**

- Number of processes completed per time unit.

- Example: Long processes may complete 1 process/hour, while short processes may complete 10 processes/second.

### 3. Turnaround Time:

- Total time taken from process submission to completion.
- Includes waiting time in the ready queue, execution time, and I/O time.
- Lower turnaround time is preferred.

### 4. Waiting Time:

- Time a process spends in the ready queue waiting for CPU allocation.
- Shorter waiting times improve overall system performance.

### 5. Response Time:

- Time from the submission of a request until the first response is produced.
- Important in interactive systems where quick feedback is required.
- Does not measure full output time, just the time to start responding.

### Summary:

These criteria are used to evaluate and compare CPU scheduling algorithms to select the most efficient one for a given system.

---

## 6. Define deadlock. What are the necessary conditions for deadlock to occur?

### Deadlock and Necessary Conditions

#### Definition of Deadlock:

Deadlock is a situation where processes are stuck in a waiting state indefinitely because the resources they need are held by other waiting processes.

#### Necessary Conditions for Deadlock:

Deadlock can occur only if all the following four conditions hold simultaneously:

#### 1. Mutual Exclusion:

- At least one resource is non-sharable, meaning only one process can use it at a time.
- If another process requests the resource, it must wait until the resource is released.

#### 2. Hold and Wait:

- A process is holding at least one resource and waiting for additional resources that are held by other processes.

#### 3. No Preemption:

- Resources cannot be forcibly taken from a process; they can only be released voluntarily after the process completes its task.

#### 4. Circular Wait:

- A circular chain of processes exists where each process is waiting for a resource held by the next process in the chain.

- For example,  $P_0$  is waiting for  $P_1$ 's resource,  $P_1$  is waiting for  $P_2$ 's resource, and  $P_n$  is waiting for  $P_0$ 's resource.

### Summary:

Deadlock occurs when these four conditions are present simultaneously. To prevent or resolve deadlocks, one or more of these conditions must be eliminated.

---

## 7. Illustrate Peterson's solution for the critical section problem.

### Peterson's Solution for the Critical Section Problem

#### Definition:

Peterson's solution is a software-based approach to ensure mutual exclusion, progress, and bounded waiting for two processes that alternate between their **critical sections** and **remainder sections**.

---

#### Key Points:

- Works for **two processes only** (labeled  $P_0$  and  $P_1$ ).
- Requires two shared variables:
  1. **int turn:**
    - Indicates whose turn it is to enter the critical section.
    - Example: If  $turn == i$ , process  $P_i$  is allowed to enter.
  2. **boolean flag[2]:**
    - Indicates whether a process is ready to enter its critical section.
    - Example: If  $flag[i] == true$ , process  $P_i$  is ready.

#### How It Works:

1. **Entering the Critical Section:**
    - Process  $P_i$ :
      - Sets  $flag[i] = true$  (shows it wants to enter).
      - Sets  $turn = j$  (allows  $P_j$  to enter if it also wants).
      - Waits until  $flag[j] == false$  or its own turn ( $turn == i$ ).
  2. **Handling Simultaneous Requests:**
    - If both processes try to enter at the same time, the turn variable determines which process enters the critical section first.
    - The process whose turn is set last will wait.
  3. **Exiting the Critical Section:**
    - After finishing,  $P_i$  sets  $flag[i] = false$  to allow  $P_j$  to enter.
-

```

do {
    flag[i] = TRUE;
    turn = j;
    while (flag[j] && turn == j)
        ; // do nothing
    critical section
    flag[i] = FALSE;

    remainder section

} while (TRUE);

```

Figure: The structure of process  $P_i$  in Peterson's solution

### Summary of Benefits:

- Ensures **mutual exclusion** (only one process in the critical section at a time).
  - Maintains **progress** (no process is unnecessarily delayed).
  - Satisfies **bounded waiting** (no process waits forever to enter).
- 

## 8. Explain different methods to recover from deadlocks.

### Methods to Recover from Deadlocks

Deadlocks can be managed in the following ways:

#### 1. Prevent or Avoid Deadlocks:

- The system uses strategies to ensure that deadlocks never occur.
- Two approaches:
  - **Deadlock Prevention:** Prevents at least one of the four necessary conditions (Mutual Exclusion, Hold and Wait, No Preemption, Circular Wait) from happening.
  - **Deadlock Avoidance:** The system is given additional information about resource usage in advance and makes decisions to avoid entering a deadlock state.

#### 2. Detect and Recover from Deadlocks:

- If a deadlock occurs, the system identifies it using algorithms and then takes steps to recover.
- Recovery can involve:
  - Terminating one or more processes to break the deadlock.
  - Preempting resources from some processes and allocating them to others.

#### 3. Ignore Deadlocks:

- The system assumes that deadlocks are rare and does nothing to prevent, detect, or recover from them.
- This approach is used by systems like UNIX, where manual intervention (e.g., restarting the system) is required if a deadlock occurs.

9. What is a semaphore? State a Dining Philosopher problem gives a solution using semaphore.

---

**10. What is demand paging? Explain the steps in handling page faults using the appropriate diagram.**

### Demand Paging

Demand paging is a memory management technique where a page is loaded into memory **only when it is needed** during program execution.

#### Key Features of Demand Paging:

- Unlike swapping entire processes, **only required pages** are swapped into memory.
- **Less I/O and memory usage:** Reduces swap time and physical memory usage.
- **Faster response time:** Unused pages are not loaded.
- **More users can be supported** as memory is used efficiently.
- Operates using a **pager** that swaps pages on demand.




Fig: Transfer of a paged memory into continuous disk space

### Handling Page Faults

A **page fault** occurs when a process tries to access a page that is not in memory. The steps to handle a page fault are:

1. **Access the Page Table:**
  - The **valid-invalid bit** in the page table is checked.
  - If the bit is valid (v), the page is in memory.
  - If the bit is invalid (i), the page is either on the disk or invalid, causing a page fault.
2. **Trigger a Page Fault Trap:**
  - The operating system (OS) detects the page fault.
3. **Determine the Page Status:**
  - Check if the page is valid but not in memory.
  - If invalid, terminate the process.



#### 4. Load the Page:

- Locate the page on the disk.
- Bring the page into memory.

#### 5. Update the Page Table:

- Mark the page as valid (v) in the page table.

#### 6. Resume Execution:

- Restart the process from where the page fault occurred.

Frame #	valid-invalid bit
	V
	V
	V
	V
	i
....	
	i
	i

page table




Fig: Page Table when some pages are not in main memory

#### Advantages of Demand Paging:

- Optimizes memory usage by loading only the required pages.
- Reduces unnecessary I/O operations.
- Supports multitasking by freeing up memory for multiple processes.

11. Discuss the structure of the page table with a suitable diagram.

### Structure of the Page Table

A **page table** maps the virtual addresses used by a program to the physical addresses in memory. There are several ways to organize the page table to improve performance and manage large address spaces.

#### 1. Hierarchical Paging

Hierarchical paging is used to handle large address spaces (32-bit or more) where the page table itself can become too large. This method divides the page table into smaller parts to reduce its size.

#### Two-Level Paging Example:

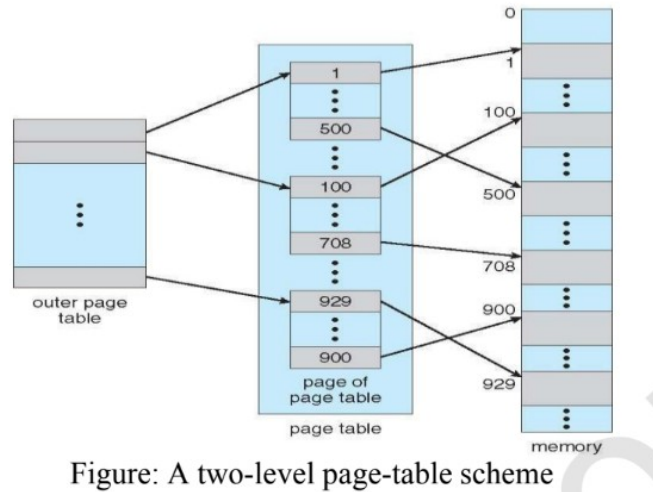



Figure: A two-level page-table scheme

Consider a system with a 32-bit logical address space and a 4 KB page size. The logical address is split into:

- **20 bits** for the page number
- **12 bits** for the page offset

Since the page table itself is large, the page number is further divided into two parts:

- **10 bits** for the first-level page table index
- **10 bits** for the second-level page table index



In this scheme, the logical address is divided into:

- $p_1$ : Index into the outer page table
- $p_2$ : Index into the inner page table

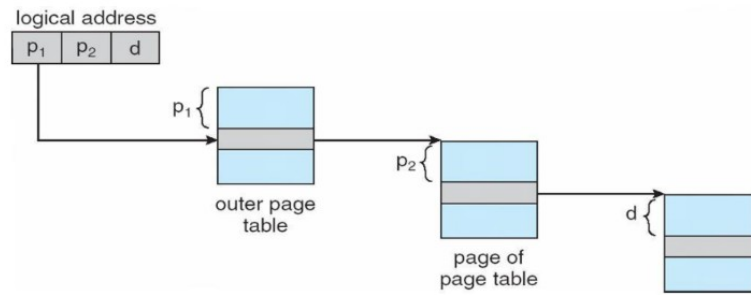


Figure: Address translation for a two-level 32-bit paging architecture

**Address translation** happens from the outer table inward, making this a **forward-mapped page table**.

## 2. Hashed Page Tables

Hashed page tables are used for larger address spaces (more than 32 bits). This method uses a hash function to map virtual page numbers to entries in a hash table.

Each hash table entry contains a **linked list** of elements that hash to the same location to handle **collisions**. Each element in the list has:

1. **Virtual page number**
2. **Mapped physical page frame**
3. **Pointer to the next element**

The process works as follows:

- The virtual page number is hashed into the table.
- The system checks the list for a match with the virtual page number.
- If a match is found, the corresponding physical address is used; if not, the system searches further down the list.




Figure: Hashed page-table

## 3. Inverted Page Tables

In inverted page tables, there is only one entry for each physical page in memory. Each entry contains:

- The **virtual address** of the page in that memory location
- Information about the **process** that owns the page

Each virtual address is represented as a triplet: <process-id, page-number, offset>. When a memory reference occurs:

1. The virtual address is checked against the inverted page table.
2. If a match is found, the physical address is generated.
3. If no match is found, it indicates an illegal memory access.




Figure: Inverted page-table

#### Advantages and Disadvantages:

- **Advantage:** It reduces the memory needed to store the page table.
  - **Disadvantages:** Searching the table is slower, and it can be challenging to implement for **shared memory** scenarios.
- 

12. Explain in detail about various file operations in a file system.

#### File Operations in a File System

A file system allows various operations on files. These operations are essential for managing files effectively and include the following:

##### 1. Creating a File

To create a file, two steps are required:

- **Space Allocation:** The system must find space for the new file in the file system.
- **Directory Entry:** An entry for the file must be added to the directory to keep track of the file's location.

##### 2. Writing to a File

To write to a file:

- A system call is made with the file name and data to be written.
- The system searches the directory to find the file's location.
- A **write pointer** keeps track of where in the file the next data should be written.
- The write pointer is updated after each write operation.

##### 3. Reading from a File

To read from a file:

- A system call is made with the file name and a location where the next block should be placed.
- The system searches the directory for the file's entry.
- A **read pointer** keeps track of where the next read will take place, and it is updated after each read.

#### 4. Repositioning within a File

This operation involves moving the file's **current file position** to a new location without performing actual I/O. This is known as **file seek**, allowing you to jump to a specific point in the file.

#### 5. Deleting a File

To delete a file:

- The system searches the directory for the file.
- Once the file is found, the file space is released, and the directory entry is erased.

#### 6. Truncating a File

Truncating a file removes its content but keeps its attributes (like name and permissions) intact. This operation resets the file's length to zero and releases its space for other uses.

#### Additional Operations:

- **Appending:** Adding new data to the end of an existing file.
- **Renaming:** Changing the name of an existing file.